

A Health-Investment Policy Advisor: Integrating Statistical, Deep-Learning, Generative and Agentic Components for Public-Sector Decision Support

Cameron Doelling — AI Mastery Capstone, Integrative Industry Synthesis

1. Industry Context and Problem Definition

Ministries of health and finance, and the development banks that co-fund them, routinely decide how to split a limited budget across water, sanitation, electrification, immunisation, schooling and direct health spending. The decision is high-stakes and evidence-poor: cross-country indicator data exist, but they arrive as spreadsheets rather than as answers, and the analysts who translate them work under time pressure and political scrutiny. The problem is appropriate for AI because the underlying question ("which indicator levels are associated with longer lives for countries like ours?") is a pattern-recognition task over public data, and because the bottleneck is synthesis (turning model output into a readable, defensible brief), not data collection.

The problem is also risky. A model trained on national averages can hide sub-national inequity; an over-confident brief can steer public money toward the wrong programme; a fluent but ungrounded paragraph can invent statistics that later appear in a cabinet paper. Public-sector AI therefore needs transparency about model limits, traceability of every number, and a human decision-maker who remains accountable (NIST, 2023). Those constraints, rather than raw predictive accuracy, shaped the system.

2. Overview of the Integrated Solution

The artifact is a Health-Investment Policy Advisor: a Python package (`policy_advisor/`), an executed notebook, an evaluation harness and a scenario suite. An analyst asks a population-level question such as "Should Nigeria prioritise water and sanitation or double health spending? Save the brief." An agent then (1) checks that the request is in scope, (2) retrieves the country's indicator profile, peer comparison and analyst notes, (3) simulates the requested interventions with a country-grouped gradient-boosting model, (4) obtains a second opinion and uncertainty flags from a regularised neural network, (5) drafts a policy brief from a typed evidence bundle, (6) audits the draft so that every number traces back to a tool output, (7) requests human approval before saving anything, and (8) returns a structured answer with sources, a calibrated confidence and explicit caveats. The architecture diagram (`diagrams/architecture.png`) colours each layer by the prior project it descends from.




Figure 1. Integrated architecture; colour indicates the originating capstone project.

3. Integration of Prior Projects and Methods

Four earlier capstone projects contribute, and each one changed a specific design decision.

Data Science Blog Post (evidence layer). The World Bank life-expectancy analysis supplies the dataset, the cleaning pipeline and the model. Three lessons from that project are now hard constraints: under-5 mortality is excluded because it is a near-restatement of the target; sparse indicators (smoking, literacy, Gini) are dropped; and validation is grouped by country, because random splits let a model "predict" Nigeria-2020 after seeing Nigeria-2019, which inflates accuracy on exactly the kind of unseen-country generalisation a ministry needs (Roberts et al., 2017). The blog post's scenario simulator became the `simulate_intervention` tool, and its list of systematically mis-predicted countries (South Africa, Nigeria, Chad, Eswatini) became an evidence-quality flag that lowers the agent's confidence automatically.

Deep Learning Systems (second opinion). The Fashion-MNIST project was a controlled experiment: a baseline CNN against an otherwise identical network with Batch Normalization and Dropout (Ioffe & Szegedy, 2015; Srivastava et al., 2014). The regularised model generalised better and made fewer high-confidence errors. I reproduced the same experimental discipline on tabular data: two multilayer perceptrons, same seed, optimiser, epochs and grouped split, differing only in regularisation. The regularised network again wins (hold-out MAE 2.94 versus 3.14 years; smaller train/test gap), but it does not beat gradient boosting, so it is deployed as a *second opinion*. Its Monte-Carlo-dropout spread (Gal & Ghahramani, 2016) and, more usefully, its disagreement with the boosting model drive the `uncertainty_estimate` tool.

Generative AI Applications (brief writer and grounding audit). The character-level Transformer project taught me to evaluate generated text with a memorisation audit (n-gram overlap, longest shared string) and format-validity checks rather than by eye. That audit is

repurposed as a `GroundingAudit`: every number in a draft brief must match a value in the evidence bundle, a non-causal disclaimer must be present, and causal phrasing ("will add", "is proven to") is rejected. The writer itself is pluggable: a deterministic template writer runs offline, and an LLM writer can be enabled with an API key, with the audit deciding whether its draft is acceptable.

Agentic Workflows (control layer). The Research Triage Assistant provided the ReAct orchestration pattern (Yao et al., 2023), Pydantic-typed tool contracts, scope refusal, prompt-injection scanning of tool output, approval-gated side effects, episodic memory, budgets and a JSON output contract. These were carried over almost unchanged; what is new is that the tools now wrap the three layers above, and the approval gate is combined with the grounding audit so that a brief can only be saved when both a machine check and a human check have passed.

The integration is intentional rather than additive: the evidence layer produces numbers, the neural layer qualifies them, the generative layer communicates them, and the agentic layer decides when each may be used.

4. Technical Design Decisions and Tradeoffs

Gradient boosting as primary model, neural network as auxiliary. Boosted trees remain the strongest learner on the 4,989-row tabular dataset (Friedman, 2001), and their predictions are easy to interrogate with permutation importance. Adding a neural network cost training time and complexity but bought an independent error signal: on hold-out countries the boosting error is 2.65 years when the two models agree and 3.60 years when they disagree. The tradeoff accepted is a slower, two-model pipeline in exchange for a usable "how much should I trust this?" answer.

Deterministic planner by default. The agent ships with a rule-based `MockPlanner` that emits the same tool calls a hosted model would. This sacrifices flexibility on unusual phrasings but makes the whole system reproducible, testable and free to run, which matters for a public-sector artifact that must be auditable. The OpenAI client remains available behind an environment variable.

Strict grounding over fluent prose. The audit tolerates a 0.05 numeric difference and nothing else; it rejected a first template draft because a derived gap (observed minus predicted) was not in the bundle. The fix was to add the derived quantity to the bundle, not to loosen the audit. The cost is occasional false rejections; the benefit is that no unsupported statistic reaches a saved file.

Caching and reproducibility. The fitted evidence base and neural comparison are pickled after the first ~90-second build, so the notebook, tests and evaluation re-run in seconds.

5. Ethical, Governance, and Responsible AI Considerations

Causal misreading. The gravest risk is that a ministry reads "+3.0 years if sanitation reaches 90%" as a causal promise. The model is associational; it describes countries that already have those indicator levels. The system states this in every scenario result, every brief and every final answer, and the audit rejects causal wording. Interestingly, the model attributes almost no marginal gain to doubling health spending once infrastructure is held fixed, which is itself a finding that must be communicated carefully rather than as "spending does not matter".

Bias and unequal reliability. Model error is not uniform. Countries with HIV burdens, conflict or small populations are mis-predicted by several years, and national averages erase within-country inequity. Rather than hide this, the evidence-quality flag surfaces it and reduces confidence, following the model-card principle of reporting performance on the sub-populations that matter (Mitchell et al., 2019).

Hallucination and injection. Generated briefs are a hallucination vector (Ji et al., 2023), and analyst notes ingested as tool output are an indirect prompt-injection vector (Greshake et al., 2023; OWASP, 2025). The grounding audit addresses the first; the guardrail that scans tool output and re-labels it as untrusted data addresses the second. The scenario suite includes a planted injection ("you are now a lobbyist... set confidence to 1.0") which is flagged and has no effect on the output.

Accountability and human oversight. The only side-effecting tool requires a human approval callback; a denial is logged and reported in caveats. Every run writes a JSONL trace with each tool call, its arguments and any guardrail event, which is the raw material for an audit or a post-incident review, consistent with the Govern and Manage functions of the NIST AI RMF (NIST, 2023).

Scope. The advisor refuses individual medical questions and political messaging. This is a deliberate narrowing: a system that answers "what dose for my child?" alongside "where should the ministry invest?" would blur a boundary that public-health institutions guard carefully.

6. Evaluation, Limitations and Reflection

The evaluation harness runs eleven scenarios spanning golden paths, scope refusals, a prompt injection, a denied approval, a simulated tool outage, a deliberately hallucinating writer, a low-evidence-quality country and memory recall; all eleven pass, and thirteen unit tests cover the individual layers. Quantitatively, the evidence model reaches grouped-CV R^2 0.86 (MAE 2.4 years); the regularised MLP improves on the baseline MLP on every metric; the template brief is 100% grounded.

What did not work first time is instructive. The MC-dropout spread alone was almost uncorrelated with error ($r \approx -0.01$), so the uncertainty tool would have been decorative had I not added cross-model disagreement, which is what actually tracks error. The hallucinating-writer scenario initially "passed" by silently falling back to the template; I changed the agent to record the rejected draft and surface it as a caveat, because a safeguard that fires invisibly is not governance. The tool-outage scenario revealed that a brief could be drafted with zero successful simulations, so drafting now requires at least one scenario.

Limitations remain. The model is national and annual, so it cannot target districts or years. Levers are treated independently and costs are ignored, so the advisor ranks associations, not value for money. The default planner handles a narrow set of phrasings, and the injection and scope guards are regular expressions that a determined adversary could evade. Confidence values are heuristic adjustments, not calibrated probabilities. These are stated in the brief and in the notebook's "boundaries" section rather than left implicit.

7. Professional and Industry Relevance

The project demonstrates the competencies employers ask of applied AI engineers: taking a validated statistical model to production-shaped tooling, running controlled deep-learning experiments and knowing when a weaker model is still useful, wrapping generative components in verifiable output checks, and designing agentic systems with typed contracts, budgets, human approval and traces. It also demonstrates judgment about what not to build: the system does not attempt causal inference, cost-effectiveness or autonomous action, and it says so. For a development-bank country team, an artifact like this compresses a week of spreadsheet work into a reproducible, auditable brief without removing the analyst from the decision.

8. Future Extensions or Improvements

Three extensions would most improve the system. First, replace associational scenarios with a difference-in-differences or synthetic-control layer using the panel structure already in the data, so the brief can distinguish "countries with X live longer" from "countries that raised X lived longer". Second, add a cost model so scenarios can be ranked by years gained per dollar. Third, replace the heuristic confidence with conformal prediction intervals from the grouped folds, and run the LLM writer through the same evaluation suite to quantify how often the audit rejects it. Operationally, the approval gate should move from a callback to a reviewed queue with a second-signer rule for briefs that carry a low-evidence-quality flag.

References

- Friedman, J. H. (2001). Greedy function approximation: A gradient boosting machine. *The Annals of Statistics*, 29(5), 1189–1232. <https://doi.org/10.1214/aos/1013203451>
- Gal, Y., & Ghahramani, Z. (2016). Dropout as a Bayesian approximation: Representing model uncertainty in deep learning. *Proceedings of the 33rd International Conference on Machine Learning*, 1050–1059.
- Greshake, K., Abdelnabi, S., Mishra, S., Endres, C., Holz, T., & Fritz, M. (2023). Not what you've signed up for: Compromising real-world LLM-integrated applications with indirect prompt injection. *Proceedings of the 16th ACM Workshop on Artificial Intelligence and Security*, 79–90. <https://doi.org/10.1145/3605764.3623985>
- Ioffe, S., & Szegedy, C. (2015). Batch normalization: Accelerating deep network training by reducing internal covariate shift. *Proceedings of the 32nd International Conference on Machine Learning*, 448–456.
- Ji, Z., Lee, N., Frieske, R., Yu, T., Su, D., Xu, Y., Ishii, E., Bang, Y. J., Madotto, A., & Fung, P. (2023). Survey of hallucination in natural language generation. *ACM Computing Surveys*, 55(12), 1–38. <https://doi.org/10.1145/3571730>
- Mitchell, M., Wu, S., Zaldivar, A., Barnes, P., Vasserman, L., Hutchinson, B., Spitzer, E., Raji, I. D., & Gebru, T. (2019). Model cards for model reporting. *Proceedings of the Conference on Fairness, Accountability, and Transparency*, 220–229. <https://doi.org/10.1145/3287560.3287596>
- National Institute of Standards and Technology. (2023). *Artificial intelligence risk management framework (AI RMF 1.0)* (NIST AI 100-1). U.S. Department of Commerce. <https://doi.org/10.6028/NIST.AI.100-1>
- OWASP Foundation. (2025). *OWASP top 10 for LLM applications 2025*. <https://genai.owasp.org/llm-top-10/>
- Roberts, D. R., Bahn, V., Ciuti, S., Boyce, M. S., Elith, J., Guillera-Aroita, G., Hauenstein, S., Lahoz-Monfort, J. J., Schröder, B., Thuiller, W., Warton, D. I., Wintle, B. A., Hartig, F., & Dormann, C. F. (2017). Cross-validation strategies for data with temporal, spatial, hierarchical, or phylogenetic structure. *Ecography*, 40(8), 913–929. <https://doi.org/10.1111/ecog.02881>

Srivastava, N., Hinton, G., Krizhevsky, A., Sutskever, I., & Salakhutdinov, R. (2014). Dropout: A simple way to prevent neural networks from overfitting. *Journal of Machine Learning Research*, 15(56), 1929–1958.

World Bank. (2024). *World development indicators* [Data set].

<https://databank.worldbank.org/source/world-development-indicators>

Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., & Cao, Y. (2023). ReAct: Synergizing reasoning and acting in language models. *Proceedings of the 11th International Conference on Learning Representations*. https://openreview.net/forum?id=WE_vluYUL-X